

When Does Graph Propagation Help Recommendation?

A Technical Review of MF-BPR and LightGCN

on Dense and Sparse Implicit Feedback

Group XX

Abstract

We review and empirically compare two representative recommender-system methods: **MF-BPR**, matrix factorization trained with the Bayesian Personalized Ranking loss, and **LightGCN**, a simplified graph convolutional network over the user-item interaction graph. Both are implemented from scratch in PyTorch and evaluated under an identical full-ranking protocol on two datasets chosen to differ mainly in interaction density: MovieLens-1M (dense, 4.8%) and Amazon Reviews 2023 Video Games (144× sparser, 0.034%). LightGCN outperforms MF-BPR on both datasets, but the margin grows by an order of magnitude with sparsity: +6.1% NDCG@10 on MovieLens-1M versus +43.5% on Video Games. Per-user analysis shows the gain concentrates on low-activity users, and a popularity analysis reveals that on sparse data part of the gain comes from recommending popular items at the cost of catalog coverage (62% → 41%). We further show a methodological pitfall: under a 300-epoch budget the layer ablation appears to exhibit “over-smoothing” beyond $L=2$, which vanishes entirely once every configuration is given a fair 600-epoch budget — deeper propagation simply converges more slowly. All results are produced by our own code and are reproducible from fixed seeds.

1 Introduction

Recommender systems learn user preferences from historical interactions. Matrix factorization (MF) has been the workhorse of collaborative filtering since the Netflix-Prize era: it embeds each user and item in a low-dimensional space and models affinity as a dot product, using only *direct* user-item interactions. A more recent line of work observes that the interaction data forms a bipartite graph, and that propagating embeddings over this graph injects *higher-order* collaborative signal — “users who liked the items I liked” — into the representation. LightGCN is the cleanest representative of this idea: it strips graph convolution down to weighted neighborhood averaging, with no feature transforms and no nonlinearities.

This project asks a focused question: **when, and for whom, does the higher-order graph signal actually help?** We keep everything else fixed — the embedding tables, the BPR training objective, the negative sampling, the evaluation protocol — so that the *only* difference between our two methods is the propagation step. Any performance gap is then attributable to the graph. We deliberately evaluate on two datasets at opposite ends of the density spectrum, because the value of borrowed neighborhood signal should depend on how much direct signal each user already has.

Our contributions are: (i) a from-scratch, ~100-line PyTorch implementation of both methods verified against published result ranges; (ii) a controlled comparison on dense and sparse data with per-user and popularity-level analyses that explain *where* the gains come from; and (iii) a cautionary methodological finding about training budgets in layer ablations.

2 Methods

2.1 Problem setup and the BPR objective

Let $\mathcal{U}$ and $\mathcal{I}$ be the user and item sets, and $\mathcal{O} = \{(u, i)\}$ the observed implicit interactions. Each user and item has an embedding $\mathbf{e}_u, \mathbf{e}_i \in \mathbb{R}^d$, and a pair is scored $\hat{y}_{ui} = \mathbf{e}_u^\top \mathbf{e}_i$. At test time we rank all items the user has not interacted with in training. Both models are trained with the pairwise BPR loss: for each observed (u, i) and a uniformly sampled unobserved item j ,

$$\mathcal{L}_{\text{BPR}} = - \sum_{(u,i,j)} \ln \sigma(\hat{y}_{ui} - \hat{y}_{uj}) + \lambda \|\Theta\|_2^2, \quad (1)$$

which pushes interacted items above random non-interacted ones in each user’s ranking.

2.2 Method 1: MF-BPR

MF-BPR uses the free embedding tables directly: $\mathbf{e}_u$ and $\mathbf{e}_i$ are the model parameters, and no other computation is performed. Its inductive bias is minimal; every user is positioned only by the items they themselves interacted with.

2.3 Method 2: LightGCN

LightGCN keeps the same tables as layer-0 embeddings $\mathbf{e}^{(0)}$ and propagates them over the interaction graph. With $\mathcal{N}_u$ the items of user u and $\mathcal{N}_i$ the users of item i ,

$$\mathbf{e}_u^{(k+1)} = \sum_{i \in \mathcal{N}_u} \frac{\mathbf{e}_i^{(k)}}{\sqrt{|\mathcal{N}_u| |\mathcal{N}_i|}}, \quad \mathbf{e}_i^{(k+1)} = \sum_{u \in \mathcal{N}_i} \frac{\mathbf{e}_u^{(k)}}{\sqrt{|\mathcal{N}_i| |\mathcal{N}_u|}}, \quad (2)$$

and the final embedding is the mean over layers, $\mathbf{e}_u = \frac{1}{L+1} \sum_{k=0}^L \mathbf{e}_u^{(k)}$. There are no weight matrices and no activation functions: LightGCN is deliberately a *simplification* of standard GCNs, which its authors showed helps rather than hurts for recommendation. Layer k mixes in k -hop neighbors, so $L=2$ already reaches other users’ items. Setting $L=0$ recovers MF-BPR exactly — the comparison is a clean ablation by construction.

2.4 Implementation

Both models are implemented from scratch in PyTorch (~100 lines total). Propagation is realized as `gather + index_add` over the edge list, which runs identically on CUDA, Apple-Silicon MPS, and CPU (sparse matmul is not supported on MPS). Because LightGCN propagates the full graph at every optimization step, we use a large batch of 65,536 with lr 3×10^{-3} , giving $9\times$ fewer propagations per epoch with convergence indistinguishable from the small-batch setting; MF uses batch 8,192 and lr 10^{-3} . Both use Adam, L_2 regularization 10^{-4} , and embedding dimension 64 unless stated. Negative sampling uses a seeded NumPy generator on the CPU, which makes runs bit-reproducible: the MF run produced *identical* test metrics on an Apple M2 and an NVIDIA L4.

3 Experimental Setup

3.1 Datasets

MovieLens-1M is the classic dense movie-rating benchmark (GroupLens Research): 1,000,209 ratings from 6,040 users. **Amazon Reviews 2023 – Video Games** (McAuley Lab, UCSD)

	MovieLens-1M	Video Games
Interactions (5-core)	999,611	814,586
Users	6,040	94,762
Items	3,416	25,612
Density (%)	4.8448	0.0336
Avg. interactions / user	165.5	8.6
Avg. interactions / item	292.6	31.8

Table 1: Datasets after preprocessing. The two datasets have interactions of the same order of magnitude but differ in density by a factor of ~ 144 .

contains 4.62M reviews; gaming involves frequent repeat purchasing, so the category retains a healthy interaction graph after filtering. Both are converted to implicit feedback (every observed pair is a positive), deduplicated, 5-core filtered, and split per user 80/10/10 into train/validation/test with seed 42 (Table ??). Figure ?? shows both are long-tailed, but at very different scales: the average MovieLens user has 165 interactions, the average Video-Games user 8.6.




Figure 1: Interactions per user (left) and per item (right), log-log.

3.2 Protocol

We evaluate by *full ranking*: for each user, all items not seen in training are scored and Recall@ K / NDCG@ K ($K \in \{10, 20\}$) are computed on the held-out items; no sampled-negative evaluation is used, as it is known to bias comparisons. Model selection uses validation NDCG@10 with early stopping (patience 30 epochs). The training budget is 300 epochs for MF (which always early-stops far below it) and 600 for LightGCN; the reason for the larger budget is discussed in Section ?? . The full queue of 18+ runs completes in ≈ 1.5 h on one NVIDIA L4 GPU.

4 Results and Analysis

4.1 Main comparison

Table ?? shows the central result. On dense MovieLens-1M, LightGCN improves NDCG@10 by +6.1%; on Video Games the same architecture change yields +43.5%. The comparison isolates the

Dataset	Model	R@10	N@10	R@20	N@20	Best ep.
MovieLens-1M	MF-BPR	0.1560	0.2456	0.2384	0.2525	45
	LightGCN ($L=3$)	0.1728	0.2605	0.2638	0.2709	450
	<i>relative gain</i>	+10.8%	+6.1%	+10.6%	+7.3%	
Video Games	MF-BPR	0.0634	0.0344	0.0969	0.0431	125
	LightGCN ($L=3$)	0.0907	0.0494	0.1349	0.0608	565
	<i>relative gain</i>	+43.2%	+43.5%	+39.3%	+41.2%	

Table 2: Test-set results at $d=64$ (R = Recall, N = NDCG). LightGCN wins everywhere, but the relative gain is $\sim 7\times$ larger on the sparse dataset.

propagation step as the only difference, so we conclude that **the value of higher-order graph signal grows sharply as direct per-user signal shrinks**. On dense data most of what the graph could tell the model is already present in each user’s own long history.

4.2 Convergence behaviour




Figure 2: Validation NDCG@10 per epoch. MF converges fast and early-stops (best epoch 45/125); LightGCN converges $\sim 10\times$ slower (best epoch 450/565) but reaches a higher plateau.

Figure ?? shows the cost of the graph: LightGCN needs both a full-graph propagation per optimization step *and* many more epochs. On Video Games the two curves cross around epoch 90 — before that, MF is the better model. In compute-constrained settings this trade-off matters as much as the final number.

The training-loss curves (Figure ??) complete the picture from the optimization side: MF fits the training pairs faster and more aggressively than LightGCN, yet generalizes worse (Table ??). The propagation step acts as a structural regularizer — each embedding is tied to its neighborhood average, so LightGCN cannot minimize the loss by memorizing individual pairs. The gap between training fit and test ranking is precisely the graph prior at work, and it also explains the slower convergence observed above.




Figure 3: Training BPR loss per epoch. MF drives the loss down faster and further (its curves stop early because of early stopping), yet ranks held-out items worse.




Figure 4: Test NDCG@10 vs. propagation layers L , all configurations trained with the same fair 600-epoch budget. Even $L=1$ clearly beats MF; no over-smoothing is visible up to $L=4$.

4.3 Ablation: propagation depth, and a training-budget pitfall

The layer count L is LightGCN’s key component: $L=0$ recovers MF, so Figure ?? directly measures the marginal value of each propagation hop. Two observations. First, a single hop already captures most of the gain on the sparse dataset (+26% NDCG@10 over MF), with diminishing returns after $L=3$. Second — and this is a finding we consider as valuable as the curve itself — our *initial* run of this ablation used a 300-epoch budget for every L and showed the classic “peak at shallow L , degradation at $L=4$ ” pattern commonly attributed to over-smoothing. Doubling the budget removed the degradation entirely: the best epoch grows from 110 at $L=1$ to ~ 565 at $L=4$, i.e. **deeper propagation converges more slowly, and an unfair training budget masquerades as over-smoothing**. Ablation tables in the literature that fix a single epoch budget across depths should be read with this caveat in mind.

4.4 Embedding dimension

Both methods improve monotonically with d (Figure ??), and the LightGCN curve is shifted upward roughly in parallel. Notably, **LightGCN at $d=16$ beats MF at $d=64$ on both datasets** despite having a quarter of the parameters: the graph prior contributes more than raw capacity does in




Figure 5: Test NDCG@10 vs. embedding dimension $d \in \{16, 32, 64\}$.

this regime.

4.5 Who benefits? Per-user analysis

Dataset	Train inter.	MF-BPR	LightGCN	Gain
MovieLens-1M	14–35	0.2877	0.3299	+14.6%
	36–75	0.2711	0.3055	+12.7%
	76–164	0.2257	0.2442	+8.2%
	≥ 165	0.1707	0.1775	+4.0%
Video Games	3–3	0.0932	0.1335	+43.3%
	4–6	0.0962	0.1348	+40.1%
	≥ 7	0.1020	0.1367	+34.0%

Table 3: Recall@20 by user-activity bucket (number of training interactions). The relative gain shrinks monotonically as activity grows.

Bucketing test users by their training degree (Table ??, Figure ??) explains where the improvement lives: on MovieLens-1M the gain falls monotonically from +14.6% for the coldest quartile to +3.9% for the most active one. Propagation lets a cold user borrow the histories of neighboring users; a heavy user’s own history already positions them well, and the borrowed signal adds little. A case study makes this concrete (notebook, Section 6.3): for a sampled 476-interaction user, MF places *seven* held-out movies in its top-10 versus LightGCN’s four — averaging in neighborhood signal can actually dilute an already well-determined profile.

4.6 Popularity bias and catalog coverage

Table ?? reveals that *how* LightGCN wins differs by regime. On dense MovieLens-1M it recommends *less* popular items than MF (mean popularity 1150→994) and covers *more* of the catalog (51%→61%): the rich graph individualizes users. On sparse Video Games the opposite happens: propagation over a thin graph concentrates mass on bestsellers (373→418) and coverage drops from 62% to 41%. Part of the +43.5% headline gain is therefore a *popularity bet*. If the product goal values discovery




Figure 6: Recall@20 per activity bucket, corresponding to Table ??.

Dataset	Model	Mean popularity	Coverage@10
MovieLens-1M	MF-BPR	1150	50.9%
	LightGCN	994	61.4%
Video Games	MF-BPR	373	62.1%
	LightGCN	418	41.4%

Table 4: Mean training popularity of top-10 recommendations and share of the catalog that ever appears in a top-10 list.

and long-tail exposure, this is a genuine weakness of LightGCN in sparse regimes — and, we believe, the most promising direction for improving the method.

5 Discussion: strengths, weaknesses, key factors

MF-BPR is fast (best epoch within 45–125), simple, trivially scalable, and surprisingly strong on dense data and heavy users; its weakness is cold users, for whom it has almost no signal. **LightGCN** delivers consistent accuracy gains that concentrate exactly where MF is weakest, and is remarkably parameter-efficient; its costs are $\sim 10\times$ slower convergence, a full-graph propagation per step, and a popularity/coverage regression on sparse graphs. The key factors affecting performance in our experiments were, in order: dataset density, training budget (especially for deeper propagation), propagation depth L , and embedding dimension. Sensible defaults from our sweeps: $L=3$, $d=64$, with the caveat that L and the epoch budget must be increased together.

6 Conclusion

We compared MF-BPR and LightGCN under a controlled protocol in which graph propagation is the only moving part. The higher-order signal helps everywhere but matters most where direct signal is scarce: sparse datasets and cold users. The gains are real but not free — slower convergence, more compute per step, and on sparse data a shift toward popular items that sacrifices 21 points of catalog coverage. Finally, our layer ablation shows that training budget can masquerade as over-smoothing, a reminder that ablations must equalize convergence, not just epochs. Future work:

a popularity-debiasing extension of LightGCN targeted at the sparse regime, which our analysis suggests is where the headroom lies.

References

- [1] S. Rendle, C. Freudenthaler, Z. Gantner, L. Schmidt-Thieme. BPR: Bayesian Personalized Ranking from Implicit Feedback. *UAI*, 2009.
- [2] X. He, K. Deng, X. Wang, Y. Li, Y. Zhang, M. Wang. LightGCN: Simplifying and Powering Graph Convolution Network for Recommendation. *SIGIR*, 2020.
- [3] X. Wang, X. He, M. Wang, F. Feng, T.-S. Chua. Neural Graph Collaborative Filtering. *SIGIR*, 2019.
- [4] F. M. Harper, J. A. Konstan. The MovieLens Datasets: History and Context. *ACM TiiS*, 2015.
- [5] Y. Hou, J. Li, Z. He, A. Yan, X. Chen, J. McAuley. Bridging Language and Items for Retrieval and Recommendation. arXiv:2403.03952, 2024.